



Featured

The AI governance mirage: Why 72% of enterprises don't have the control and security they think they do

Matt Marshall

12:04 pm, PT, April 21, 2026



Decision makers at **72% of organizations** claim to have two or more AI platforms that they identify as their "primary" layer, according to a survey of 40 enterprise companies conducted by VentureBeat last month, revealing real gaps in security and control.


For enterprise management and technical leaders, and especially security leaders, these multiple AI platforms extend the attack surfaces of most enterprises at a time when AI-driven attacks have become increasingly potent.

Skip Ad



0:05 / 0:12





Source: VentureBeat Pulse Research Q1 2026

The multiple platforms — which include offerings from hyperscaler or AI labs like Microsoft Azure, Google, OpenAI or Anthropic, or big application companies like Epic, Workday or ServiceNow — reflect a state of sprawl that has emerged as these big software providers rush to offer their own AI to their enterprise customers.

Those customers, in their own rush to scale AI, are finding they aren't building a singular strategy — in fact they may be building a collection of contradictions.

The strategic paradox: why leading enterprises are building around their vendors

For example, take the strategic paradox faced by Mass General Brigham (MGB) hospital system, which has 90,000 employees and is the largest employer in Massachusetts. The hospital system last year had to shut down an uncontrolled number of internal proof of concepts that had sprouted up as employees had gotten carried away with AI projects, said CTO Nallan “Sri” Sriraman at the [VentureBeat AI Impact](#) event in Boston on March 26, which focused on the challenges of scaling AI.

Instead, the company decided it was better to wait for the software giants it already uses to deliver on their AI roadmaps. Since these companies have so many resources, and were making AI a top priority themselves, it made no sense for MGB to try to build its own AI layer that would be duplicative, he said. “Why are we building it ourselves?” he asked. “Leverage it.”

Yet, even then, Sriraman’s team has been forced to build workarounds, where those companies haven’t done enough.

For example, MGB has just completed a “full-scaled” custom build around Microsoft’s Copilot — to get essentially everything offered by that tool — by putting a “skin” around Copilot to handle the safety and data privacy concerns the major model providers haven’t yet mastered. Specifically, MGB needed a way for employees to prompt the AI and not have their protected health information (PHI) leaked back to the Copilot LLM provider, OpenAI. The new secure platform, which can support up to 30,000 users, is really the ultimate contradiction: Even though the company has a mandate to leverage the AI provided by the bigger companies, it needs to build around its failures.

The contradiction goes even further. These software vendors used by MGB — which also include Epic, Workday and ServiceNow — are all now building agents for their AI, all operating differently. So MGB has to invest in building a “control plane that coordinates and orchestrates all of these agents,” Sriraman said. “That’s where our investment is going to be.”


He noted that companies like his are “discovering and experimenting as the landscape keeps shifting.” The marketplace is “still nascent,” he said, which makes decisions difficult.

The "six blind men" problem

Sriraman explained the current vendor landscape with an analogy: “When you ask six blind men to touch an elephant and say, what does this elephant look like?” Sriraman said. “You’re gonna get six different answers.”

What emerges from the research VentureBeat conducted in the first quarter, along with conversations like the one in Boston, is a situation that we at VentureBeat are calling a “governance mirage.” While many enterprises say they have adequate governance, in reality they haven’t created clear accountability or specific guardrails, evaluations or security processes to ensure that governance.

The data of disconnect: confidence vs. systematic oversight



Source: VentureBeat Pulse Research Q1 2026

The research comes from surveys across January, February and March by VentureBeat of enterprise companies with 100 or more employees, with 40 to 70 qualified respondents per topic area — covering agentic orchestration, AI security, RAG and governance. The data lacks statistical significance in many areas and should be treated as directional.

The research on governance found that a majority, or 56%, of respondents said they are “very confident” that they’d detect a misbehaving AI model, suggesting that most decision-makers believe they have sufficient basic governance at their companies.

However, nearly a third of respondents have no systematic mechanism to detect AI misbehavior until it surfaces through users or audits. In a world where telemetry leakage accounts for 34% of GenAI incidents ([Wiz](#)), and the global average breach cost has hit \$4.4M ([IBM 2025 Cost of a Data Breach](#)), finding out after the damage is done is the default for too many companies.

Moreover, 43% of respondents say a central team owns AI governance. That sounds reassuring — until you look at what’s happening everywhere else. Twenty-three percent say governance is unclear or actively contested between teams. Twenty percent say each platform team governs independently. Six percent say no one has formally addressed it. The rest said they were unsure who owned it.

More telling is the barrier data. When asked about the single biggest obstacle to governing AI across platforms, “no single owner or accountable team” ranked second at 29% — just behind vendor opacity. Accountability structure and lack of vendor transparency are the two dominant failure modes, and they compound each other: Without a central owner, no one has the mandate to demand transparency from the vendors.

The day-two bill: managing sprawl, creep, and lock-in

The scaling trap: Red Hat's warning

Brian Gracely, Senior Director at Red Hat, who also spoke at the VentureBeat Boston event last month, addressed the infrastructure side of this sprawl, warning that many enterprises are falling into a trap of deceptive initial wins.

Gracely noted that the barrier to entry is almost nonexistent at the start, with nearly anyone able to spin up a project using a credit card and an API key. "Day zero is very, very easy," Gracely said. "Day two is when the bill comes due."

Red Hat is positioning its software layer (OpenShift AI) as the necessary buffer to prevent enterprises from getting buried in a single provider's proprietary ecosystem. Gracely's point is direct: If your control system is built entirely inside one cloud provider's toolset, you are effectively "renting a cage." The illusion of speed in the early pilot phase often hides a technical debt that becomes obvious the moment you try to move your AI work to a different platform.

Gracely illustrated this with a recent example. A senior leader from Red Hat's centralized CTO office spent part of her vacation contributing to an open-source agent project called OpenClaw, which became widely popular in the first quarter. Within days of her name appearing as a project maintainer, Red Hat was fielding calls from major New York banks. Their problem was immediate: They realized they already had upwards of 10,000 employees bringing "claws" — agent-based tools — into their infrastructure with zero centralized oversight.

Breaches caused by employees working on these sorts of unapproved technologies are costly. These so-called "[shadow AI](#)" incidents cost on average [\\$670K more](#) than standard incidents, according to IBM.

Red Hat's Gracely noted that while organizations can try to shut down these unapproved ports, they eventually have to figure out how to make them productive and secure — a task that requires a serious investment in an orchestration or platform layer.

The dynamic defensive: MassMutual's refusal to bet


While some enterprise companies seek an "[AI operating system](#)" that oversees all of their AI technologies and apps, others are simply refusing to sign the check. Sears Merritt, CIO and head of enterprise technology at MassMutual, is managing the governance conundrum by intentionally staying in a state of high-velocity flexibility.

"Things are so dynamic, it's hard to know which of the AI vendors will end up on top," Merritt said at the Boston event. For that reason, MassMutual is refusing to enter any long-term contracts with AI vendors. Merritt's strategy of "dynamic defensive" highlights a core finding of our research: Vendor popularity is changing radically month to month.

Anthropic, for example, went from 0% in January to nearly 6% in February, in the number of respondents reporting what agent orchestration technology they were using. Again, the sample size was small, at 70 respondents. Still, even if directional, the dynamic landscape suggests picking a "primary" winner today is a fool's errand.

THE ANTHROPIC JUGGERNAUT

Enterprise Adoption Growth: January Baseline vs. February Surge



Source: VentureBeat Pulse Research Q1 2026

The January figure likely reflects survey composition: Respondents represent the broader enterprise market, not the developer community where Anthropic has seen its strongest early traction.

Until recently, most organizations had signed up early with leaders like Microsoft and OpenAI as their main orchestration providers, due to their early lead with Copilot. Our finding that Anthropic is just now pushing into enterprise agent orchestration may be a confirmation of the recent excitement around that platform.

One possible explanation is that enterprises already using Claude for model inference are now routing through Anthropic's native tooling rather than third-party frameworks — though the sample is too small to draw firm conclusions.

The rise of "platform creep"

The leading providers are also shifting toward "managed agents," as reflected by [Anthropic's recent announcement](#). This offering suggests possible continued platform creep, whereby providers like OpenAI


and Anthropic [take over more and more of the AI infrastructure](#) — most specifically, in this case, the memory of agentic session details. And there the trap is set. Once your session data and orchestration live inside a provider's proprietary database, you aren't just using a model; you are living in its ecosystem.

Moreover, persistent agent memory is a prime target for memory poisoning via injected instructions that influence every future interaction. And when that memory lives in a provider's database, you lose your own forensic capability.

The security irony: The fox guarding the hen house

We are seeing this platform creep in our data as well. The most jarring finding in our Q1 data is what we call the "Security Irony": the fact that the providers most responsible for creating enterprise AI risk are the same ones enterprises are using to manage it.

Respondents said the top selection criterion for AI orchestration platforms was "security and permissions generally" (37.1%), beating out other criteria like cost, flexibility, control and ease of development. Yet, the market is choosing convenience over sovereignty. According to our survey, 26% of enterprises in February were using OpenAI as their primary security solution — the very same provider whose models create the risks they are trying to secure. That trend only seemed to strengthen in March, though, as stated before, we want to be careful. Our sample size is small, and this data should only be taken as directional.



Source: VentureBeat Pulse Research Q1 2026

It's not clear whether enterprises are choosing OpenAI as a security solution, or just relying on its built-in security features offered by Microsoft Azure (which partnered with OpenAI when it pushed its Copilot solution aggressively in 2024) because customers were already on that platform.

Beyond the data, there are anecdotal signs that OpenAI's enterprise position may be shifting. Anthropic's Claude Code drew significant attention among developers early this year alongside the Claude 4.6 model. The subsequent announcement of [Mythos, its security-focused model](#), prompted interest from enterprise security teams given its ability to [identify vulnerabilities](#). OpenAI has [also announced a security-focused model](#), GPT-5.4-Cyber.

Our data may also point to a drop in OpenAI's relative position in a few enterprise AI categories. One area was data-retrieval, where OpenAI again leads among third-party providers, but we saw an increase in the number of respondents instead using in-house solutions for retrieval — perhaps a sign that AI models and agents are getting better at natively being able to use tools to call directly to companies' existing databases, and that custom code is often a way companies are building this in. However, here again we feel our data is at best directional for now.

We are asking the fox to guard the hen house. Hyperscaler security features (like those from OpenAI, Azure, and Google) are winning, because they are already integrated into the platforms enterprises are using. But it creates a single-provider dependency. As agents gain the power to modify documents, call APIs and access databases, the "governance mirage" suggests we have control, while the data shows we are simply clicking "I agree" on whatever the hyperscalers offer. The resulting risks, however, include content injection, privilege escalation and data exfiltration.

The path forward: toward a unified control plane

The search for the "Dynatrace for AI"

So, what is the way out? Sriraman argued that the industry desperately needs a "central observability platform" — a "Dynatrace for AI" — that provides full end-to-end visibility, including model drift and safety prompting, agent behavior analytics, privilege escalation alerts, and forensic logging. He is currently working with a number of potential providers to deliver on this.

The "swivel chair" warning

Sriraman warned that without a unified control plane, enterprises are at risk of sliding back into a fragmented "swivel chair" world — reminiscent of the early, inefficient days of Robotic Process Automation (RPA) — where employees are forced to constantly jump between different siloed AI tools to finish a single workflow. "We don't want to create a world where you have to switch to do something here and then go back to the platform to do something else," he said.

But that desire for a single control plane conflicts with the desire to avoid lock-in. Our data shows the market has settled on the "hybrid control plane." In other words, the most popular situation among our respondents (at 34.3%), was to use model provider-native solutions like Copilot Studio or OpenAI assistants for some workflows, while also running external options like LangGraph or custom orchestration for others. Smaller numbers of companies reported being more dogmatic here, whether that be deliberately removing the model provider from the orchestration layer entirely, relying only on custom orchestration tools, or relying only on the model provider's technology

Enterprises trust no single provider enough to give them full control, yet they lack the engineering capacity to build entirely from scratch.

The bottom line: The “big red button”

Visibility and integration are only half the battle. In a high-stakes industry like healthcare, Sriraman argues that any legitimate control plane must also offer a hard-stop capability. “We need a big red button,” he said. “Kill it. We should be able to have that ... without that, don’t put anything in the operational setting.” In fact, such a kill switch was formally [called for by the security community group OWASP](#) as part of a [recommended security framework](#).

The “governance mirage” is the belief that you can scale AI without deciding who owns the control and security plane.

If you are one of the 72% of organizations claiming multiple “primary” platforms, be careful because you may not have a strategy; you may have a conflict of interest. It suggests that the winner of the war between the AI behemoths — OpenAI, Anthropic, Google, Microsoft, etc. — won’t necessarily be the one with the best model, but the one that manages to sit *above* the models and help enterprises enforce a single version of the truth. That may be difficult to achieve, though, given that companies won’t want lock-in with a single player.

The data suggests enterprises are already resisting that outcome — and may need to formalize that resistance. Enterprises arguably need to own their control plane with independent security instrumentation, not wait for a vendor to win that role for them.

Subscribe to get latest news!

Deep insights for enterprise AI, data, and security leaders

VB Daily

AI Weekly

AGI Weekly

Security Weekly

Data Infrastructure Weekly

VB Events

All of them

Enter Your Email

By submitting your email, you agree to our [Terms](#) and [Privacy Notice](#).

Get updates

More




Google's new Deep Research and Deep Research Max agents can search the web and your private data

The release, built on Google's Gemini 3.1 Pro model, marks an inflection point in the rapidly intensifying race to build AI systems that can autonomously conduct the kind of exhaustive, multi-source research that has traditionally consumed hours or days of human analyst time. It also represents Google's clearest bid y...

Michael Nuñez

April 21, 2026



Kimi K2.6 runs agents for days — and exposes the limits of enterprise orchestration

Agents running for hours and days are breaking enterprise orchestration. Kimi K2.6 is making the case that open-source models can manage what current frameworks can't.

Emilia David

April 21, 2026




What AI model should you use for revenue intelligence? Von says all the big ones, and it will automate mixing and matching for you

The company is currently backed by several tier-one venture capital firms, including Sequoia Capital, Lightspeed, Insight Partners, and GV (Google Ventures).

Carl Franzen

April 21, 2026



Train-to-Test scaling explained: How to optimize your end-to-end AI compute budget for inference

AI reasoning does not necessarily require spending huge amounts on frontier models. Instead, smaller models can yield stronger performance on complex tasks while keeping per-query inference costs manageable

Ben Dickson

April 17, 2026



Anthropic just launched Claude Design, an AI tool that turns prompts into prototypes and challenges Figma

The simultaneous launches mark a watershed for Anthropic, whose ambitions now visibly extend from foundation model provider to full-stack product company — one that wants to own the arc from a rough idea to a shipped product. The timing is also significant: Anthropic hit roughly \$20 billion in annualized...

Michael Nuñez

April 17, 2026



Should my enterprise AI agent do that? NanoClaw and Vercel launch easier agentic policy setting and approval dialogs across 15 messaging apps

By integrating Vercel's Chat SDK and OneCLI's credential vault, NanoClaw 2.0 ensures that no sensitive action occurs without explicit human consent, delivered natively through the messaging apps where users already live.

Carl Franzen

April 17, 2026



Salesforce launches Headless 360 to turn its entire platform into infrastructure for AI agents

The announcement, made at the company's annual TDX developer conference in San Francisco, ships more than 100 new tools and skills immediately available to developers. It marks a decisive response to the existential question hanging over enterprise software: In a world where AI agents can reason, plan, and...

Michael Nuñez

April 16, 2026



Meta researchers introduce 'hyperagents' to unlock self-improving AI for non-coding tasks

Creating self-improving AI systems is an important step toward deploying agents in dynamic environments, especially in enterprise production environments, where tasks are not always predictable, nor consistent.

Ben Dickson

April 15, 2026



We tested Anthropic's redesigned Claude Code desktop app and 'Routines' — here's what enterprises should know

For the enterprise, the Desktop GUI is likely to become the standard for management and review, while the CLI remains the tool for execution.

Carl Franzen

April 15, 2026



AI's next bottleneck isn't the models — it's whether agents can think together

Cisco's Vijoy Pandey on the infrastructure gap holding back next-gen AI: agents can be stitched together, but there's no shared context or intent.

Taryn Plumb

April 15, 2026



Traza raises \$2.1 million led by Base10 to automate procurement workflows with AI

For decades, procurement has been the back office that enterprise software forgot. Billions of dollars flow through vendor negotiations, purchase orders, and supplier communications every year at the largest manufacturers and construction companies in the country — and the vast majority of that work still runs...

Michael Nuñez

April 15, 2026



Anthropic's Claude Managed Agents gives enterprises a new one-stop shop but raises vendor 'lock-in' risk

Anthropic announced a new platform last week, Claude Managed Agents, aiming to cut out the more complex parts of AI agent deployment for enterprises and competes with existing orchestration frameworks.

Emilia David

April 14, 2026



Press Releases

[Contact Us](#)

[Advertise](#)

[Share a News Tip](#)

[Contribute](#)

[Privacy Policy](#)

[Terms of Service](#)

[Consent Preferences](#)

[Do Not Sell or Share My Personal Information](#)

[Limit the Use Of My Sensitive Personal Information](#)

© 2026 VentureBeat. All rights reserved.