



math systems. When a coding harness outperforms a domain-specific agent, you've learned something about harnesses, not math. The generalization is the finding.

- **The verifiability spectrum.** A framework for classifying your work into machine-checkable, expert-checkable, and genuinely judgment-dependent tiers — and why the judgment-dependent bucket is smaller than you think.
- **Why architectural advances beat the model curve.** Organizational insights transfer across domains at near-zero cost. That's the mechanism behind the METR data showing task completion horizons doubling every four to seven months.
- **What “sniff-checking” actually means for your career.** The evaluation meta-skills in your domain that are about to get more valuable, not less — and the ones that are about to get cheaper.

But first — how a tool built to write code ended up producing better mathematics than the mathematicians' own solution, and what that means for the way you work right now.

LINK: Grab the prompts

After running the kit, you'll have a concrete map of which parts of your job are ready for structured AI delegation today — and which ones need verification infrastructure before you hand them off. You'll know where you've been overestimating how much of your work is genuinely judgment-dependent versus “involved but routine.” The second prompt gives you





math systems. When a coding harness outperforms a domain-specific agent, you've learned something about harnesses, not math. The generalization is the finding.

- **The verifiability spectrum.** A framework for classifying your work into machine-checkable, expert-checkable, and genuinely judgment-dependent tiers — and why the judgment-dependent bucket is smaller than you think.
- **Why architectural advances beat the model curve.** Organizational insights transfer across domains at near-zero cost. That's the mechanism behind the METR data showing task completion horizons doubling every four to seven months.
- **What “sniff-checking” actually means for your career.** The evaluation meta-skills in your domain that are about to get more valuable, not less — and the ones that are about to get cheaper.

But first — how a tool built to write code ended up producing better mathematics than the mathematicians' own solution, and what that means for the way you work right now.

LINK: Grab the prompts

After running the kit, you'll have a concrete map of which parts of your job are ready for structured AI delegation today — and which ones need verification infrastructure before you hand them off. You'll know where you've been overestimating how much of your work is genuinely judgment-dependent versus “involved but routine.” The second prompt gives you





math systems. When a coding harness outperforms a domain-specific agent, you've learned something about harnesses, not math. The generalization is the finding.

- **The verifiability spectrum.** A framework for classifying your work into machine-checkable, expert-checkable, and genuinely judgment-dependent tiers — and why the judgment-dependent bucket is smaller than you think.
- **Why architectural advances beat the model curve.** Organizational insights transfer across domains at near-zero cost. That's the mechanism behind the METR data showing task completion horizons doubling every four to seven months.
- **What “sniff-checking” actually means for your career.** The evaluation meta-skills in your domain that are about to get more valuable, not less — and the ones that are about to get cheaper.

But first — how a tool built to write code ended up producing better mathematics than the mathematicians' own solution, and what that means for the way you work right now.

LINK: Grab the prompts

After running the kit, you'll have a concrete map of which parts of your job are ready for structured AI delegation today — and which ones need verification infrastructure before you hand them off. You'll know where you've been overestimating how much of your work is genuinely judgment-dependent versus “involved but routine.” The second prompt gives you





math systems. When a coding harness outperforms a domain-specific agent, you've learned something about harnesses, not math. The generalization is the finding.

- **The verifiability spectrum.** A framework for classifying your work into machine-checkable, expert-checkable, and genuinely judgment-dependent tiers — and why the judgment-dependent bucket is smaller than you think.
- **Why architectural advances beat the model curve.** Organizational insights transfer across domains at near-zero cost. That's the mechanism behind the METR data showing task completion horizons doubling every four to seven months.
- **What “sniff-checking” actually means for your career.** The evaluation meta-skills in your domain that are about to get more valuable, not less — and the ones that are about to get cheaper.

But first — how a tool built to write code ended up producing better mathematics than the mathematicians' own solution, and what that means for the way you work right now.

LINK: Grab the prompts

After running the kit, you'll have a concrete map of which parts of your job are ready for structured AI delegation today — and which ones need verification infrastructure before you hand them off. You'll know where you've been overestimating how much of your work is genuinely judgment-dependent versus “involved but routine.” The second prompt gives you





math systems. When a coding harness outperforms a domain-specific agent, you've learned something about harnesses, not math. The generalization is the finding.

- **The verifiability spectrum.** A framework for classifying your work into machine-checkable, expert-checkable, and genuinely judgment-dependent tiers — and why the judgment-dependent bucket is smaller than you think.
- **Why architectural advances beat the model curve.** Organizational insights transfer across domains at near-zero cost. That's the mechanism behind the METR data showing task completion horizons doubling every four to seven months.
- **What “sniff-checking” actually means for your career.** The evaluation meta-skills in your domain that are about to get more valuable, not less — and the ones that are about to get cheaper.

But first — how a tool built to write code ended up producing better mathematics than the mathematicians' own solution, and what that means for the way you work right now.

LINK: Grab the prompts

After running the kit, you'll have a concrete map of which parts of your job are ready for structured AI delegation today — and which ones need verification infrastructure before you hand them off. You'll know where you've been overestimating how much of your work is genuinely judgment-dependent versus “involved but routine.” The second prompt gives you

