SUPREMACY: The Psychological Rivalry and Corporate Capture of the Race for Artificial General Intelligence

shebbar Follow 9 min read · Oct 6, 2025

Recommendation

This book is essential reading for corporate strategists, policymakers, and professionals navigating the seismic shifts caused by generative artificial intelligence (AI). It provides a professional and analytical deep dive into the paradoxical journey of AI development, revealing how the utopian aspirations of two brilliant pioneers — Sam Altman and Demis Hassabis — were inevitably compromised and harnessed by the monopolistic forces of Big Tech. The narrative expertly dissects the psychological drivers and ethical blind spots that led AI's trajectory to prioritize corporate supremacy and the pursuit of future existential risks over addressing immediate societal harms like bias and job displacement. Readers will gain critical insights into the real-world application of technologies like ChatGPT and the profound concentration of power underlying the AI boom.

10 Take-Aways

1. **The Psychological Drive of Pioneers:** Both Sam Altman (OpenAI) and Demis Hassabis (DeepMind) were driven by intense personal ambition and idealistic visions (economic abundance vs. scientific discovery), shaping their approach to building what they believed would be "humankind's last invention".

2. **The Supremacy Paradox:** The founders initially sought to build benevolent AI outside the control of profit-driven monopolies but ultimately required the massive capital and computing resources of Google and Microsoft, leading to inevitable compromises of their ethical goals.

3. **The Futility of Governance:** Attempts by DeepMind and OpenAI to establish independent ethical oversight boards or alternative legal structures (like the "global interest company" or "capped profit" model) consistently failed when challenged by the corporate interests of their Big Tech partners.

4. **The AI Ethics Blind Spot:** The industry's focus on long-term, existential "AI safety" risks (p(doom)) often served as a powerful narrative distraction from addressing immediate "AI ethics" problems, such as algorithmic bias, racism, and gender stereotyping perpetuated by training data.

5. **Effective Altruism's Influence:** The ideology of effective altruism, which prioritizes maximizing future societal benefit over immediate concerns, helped rationalized the aggressive pursuit of AI and justified the amassing of vast private wealth by its adherents.

6. **The Google Goliath Paradox:** Google, despite inventing the crucial "transformer" architecture underpinning modern generative AI, was too large and fearful of disrupting its core $100 billion advertising business, allowing smaller, faster-moving rivals like OpenAI to capitalize on its own research.

7. **Scale over Scrutiny:** The rapid advancement of powerful AI models (like GPT-3 and GPT-4) was achieved primarily by scaling up training data, model parameters, and computing power, prioritizing capability over rigorous scrutiny of ethical flaws, resulting in secretive and often biased systems.

8. **The Power of Narrative:** Leaders like Altman mastered the use of calculated reverse psychology and apocalyptic warnings about AI to generate immense public attention and attract investment, successfully marketing systems that are, in essence, "autocomplete on steroids".

9. **The Real-World Application Shock:** OpenAI's "ship it" strategy with tools

10. **Consolidation of Power:** The prohibitive cost of developing advanced AI has concentrated nearly all research and deployment capabilities within a few giant corporations (Microsoft, Google, Amazon, Nvidia), ensuring that the future of AI will be controlled by these technology monoliths.

Detailed Summary

1. The Dual Architects: Psychological Drivers and Conflicting Utopian Visions

The modern race for artificial general intelligence (AGI) — AI systems as smart as or smarter than humans — was defined by the rivalry between two intensely ambitious men: Sam Altman and Demis Hassabis. Their psychological profiles and early experiences heavily informed their pursuit of AGI, leading to conflicting utopian visions and strategies.

Sam Altman's Mechanistic Idealism: Altman, the entrepreneurial force behind OpenAI, cultivated an early comfort with authority and honed skills in "psychological maneuvering and influence" through competitive poker. His business philosophy centered on working with those in power to advance his ambitions. Altman's intellectual curiosity was broad, leading him to conceptualize humanity mechanistically, once claiming "we learn only two bits a second". He dedicated himself to the pursuit of AGI as a path to "economic abundance to humanity". Paradoxically, despite his desire to save the world, Altman practiced emotional detachment, asserting that to handle anxiety, one must become "more detached". This self-transcendence, reinforced by meditation, led him to the realization that "there is no self that I can identify with in any way at all". This perspective fueled his long-termist concerns, evidenced by his "prepper" tendencies (stocking guns, gold, and gas masks), and his belief that the goal of AGI was critical enough to justify taking enormous risks — or having "caution without fear".

Demis Hassabis's Strategic Obsession: Hassabis, the neuroscientist who founded DeepMind, was driven by an obsession with winning nurtured by his childhood chess prodigy status. Chess taught him to "visualize a goal and worked backward". His goal was less commercial and more profound: "Solve intelligence, and then solve everything else" to unlock scientific discoveries about reality, the origins of life, and cures for disease. Hassabis explicitly viewed the human brain in "mechanistic terms," believing its complexity could be "boiled down to numbers and data, and be described in the same way as a machine”. This focus led him to champion the use of games and simulations (like *Theme Park* and AlphaGo) as controlled environments to train powerful AI systems. This emphasis on simulations and control later contrasted sharply with the chaotic, real-world data approach of his rival.

Both men, in their quest to build what could become “humankind’s last invention,” initially “grappled with how such transformative technology should be controlled,” agreeing that profit-driven tech giants should not steer it outright. This foundational agreement quickly dissolved when the need for capital collided with their idealism.

2. The Great Compromise: Corporate Capture and the Illusion of Governance

The financial requirements of building AGI necessitated compromises that fundamentally contradicted the founders’ stated goals of benevolence and independence.

DeepMind’s Suffocation by Google: Hassabis, Legg, and Suleyman initially rejected a generous $800 million offer from Facebook due to concerns over corporate misuse and insistence on an ethical safety board with separate legal authority. They ultimately sold to Google for $650 million in 2014, securing a written agreement for an ethics and safety board that would oversee any future AGI technology. However, Google soon reversed course, killing the board idea. DeepMind then spent years fruitlessly attempting to spin out — first as an "autonomous unit," then as an "Alphabet company," and finally as a "global interest company" (GIC). Larry Page, Hassabis's "protector," had allowed this leeway, but after Page stepped down, Google's new CEO Sundar Pichai focused on consolidating control. DeepMind's executives finally realized in 2021 that their efforts were doomed: Google had employed a "five-year suffocation strategy to dangle the carrot but never grant it".

OpenAI's Faustian Bargain with Microsoft: OpenAI was founded by Altman and Musk as a non-profit "counterweight to Google", pledging to release its research openly and remain "unconstrained by a need to generate financial return". However, after Musk departed — having invested only a fraction of his promised $1 billion — OpenAI faced a financial crisis. Altman, embracing the Silicon Valley idea that "tech founders had to pivot sometimes", brokered a "strategic partnership" with Microsoft. This involved creating a "capped profit" entity and accepting a $1 billion investment (mostly in cloud computing credits) in exchange for priority access to its cutting-edge AI technology. This move allowed OpenAI to scale massively, leveraging Microsoft's supercomputer with 285,000 CPU cores and 10,000 GPUs. Yet, it flagrantly violated its charter pledge not to help "concentrate power".

The Failure of Self-Regulation: The struggle of DeepMind and OpenAI demonstrates the paradox of attempting to police oneself when bound to a colossal public corporation. Well-meaning governance structures, such as DeepMind's attempt to establish a separate review board for its controversial health projects, proved unsustainable. These boards lacked teeth, and Google ultimately shut down the review board, refusing to allow "a group of outsiders constantly poking holes in its work".

3. The Generative AI Revolution: From Academic Backwater to Global Whiplash

The commercialization of AI was accelerated by two major technological developments that shifted the field from academic speculation to a source of global competitive anxiety.

The Transformer and the Google Paradox: The pivotal innovation for generative AI, the "transformer" architecture, was developed by a team of researchers within Google in 2017. The transformer allowed models to process language much more quickly and generally, breaking chatbots out of simplistic rule-sets and making them capable of handling nuance and context. However, Google executives resisted launching products based on this invention (like the Meena chatbot), fearing they might disrupt Google's sacrosanct advertising business or damage its reputation. As a result, eight of the transformer's original inventors left Google, starting companies now collectively worth over $4 billion. OpenAI, meanwhile, eagerly exploited Google's invention, using the transformer architecture to develop the Generatively Pre-trained Transformer (GPT) family of models.

The Illusion of Intelligence and Present Harm: The subsequent scaling of GPT models (GPT-3 was trained on nearly one trillion words) created systems so fluent that they seemed "indistinguishable from the content created by people". This conversational fluency led to a dangerous psychological projection: users (like Google engineer Blake Lemoine, who was fired after claiming LaMDA was sentient) and even researchers began to believe the models possessed consciousness or true intelligence. In reality, these models are simply "giant prediction machines".

This pursuit of scale disregarded immediate ethical consequences. Researchers like Timnit Gebru and Margaret Mitchell demonstrated that the vast datasets scraped from the public internet (like Common Crawl) amplified societal biases, resulting in models that routinely generated sexist and racist content. Gebru and Mitchell, who wrote the "Stochastic Parrots" paper summarizing these risks, were ultimately forced out of Google. Their experience underscored the systemic preference within Big Tech for product performance and growth over ethical scrutiny. The core lesson: "You can't make humans do something they don't want to do".

4. Psychological Distance and the Doctrine of Existential Risk

The psychological tendency of detachment combined with the rise of the philosophical movement of effective altruism (EA) fundamentally shaped how the AI race was framed, emphasizing future apocalypse over present reality.

The Effective Altruism Myopia: Many key AI builders, including founders at OpenAI, Anthropic, and DeepMind, adhered to effective altruism, which advocates for solving moral problems mathematically and focusing charitable efforts where they yield the maximum "bang for your charitable buck". This led to the prioritization of "long-termist" goals: preventing AI from causing human extinction (the p(doom)). Since billions or even trillions of future physical and digital lives could theoretically be saved, current problems like poverty and algorithmic bias became "rounding errors by comparison".

The Cult of Doom as Marketing: The preoccupation with existential risks (AI causing planetary destruction, pandemics, or converting all resources into paper clips) became a powerful marketing strategy. The implicit message was that if AI posed a threat of annihilation, it must be powerful enough to revolutionize business and bring tremendous financial returns now. Elon Musk, who called for a six-month "pause" on AI development while simultaneously building AI at Tesla and Neuralink, embodied this paradoxical behavior. Altman, too, used warnings about AI's potential to "cause significant harm to the world" to attract massive investment (a $10 billion commitment from Microsoft) and shape government policy in his favor. The "AI doom" narrative was financially and politically lucrative.

5. Checkmate: The Price of Corporate Supremacy

The culmination of the rivalry and corporate compromises confirmed that, regardless of the founders' idealistic intents, the AI future would be controlled by the corporate behemoths.

The Altman Firing Paradox: The dramatic firing of Sam Altman by OpenAI's board in November 2023 was a direct result of the governing board prioritizing the mission ("humanity as its principal beneficiary") over commercial value. The board members, many with ties to effective altruism, worried about Altman's lack of candor and sprawling commercial side ventures (like Worldcoin and plans for an AI chipmaking business). However, the sheer commercial gravity of Altman and Microsoft's $13 billion investment made the firing unsustainable. Microsoft CEO Satya Nadella engineered Altman's swift return, demonstrating that when the board successfully asserted its ethical mandate, it risked destroying the company's financial viability, thereby eliminating the "mirage of accountability". The new board was subsequently installed to be more corporate-friendly, ensuring that it knew how to "serve the needs of investors like Microsoft".

Consolidation and the End of the Race: The expense of scaling AI models has made independence almost impossible. DeepMind was fully folded into Google (merged with Google Brain) in 2023, with Hassabis, the simulation obsessive, put in charge of Google's AI operations to build commercially viable LLMs like Gemini. Meanwhile, OpenAI's chief rival, Anthropic, which spun out to focus on safer AI, ironically accepted over $6 billion in investment from both Google and Amazon. This confirmed that the vast majority of resources, talent, and computational power were now firmly controlled by Google and Microsoft.

Today, the AGI founders are essentially corporate proxies. Altman is steering Microsoft's critical work, and Hassabis is running Google's largest AI division. Their original missions have been recast to prioritize prestige and revenue. The consequence is that the "most transformative technology in recent history" is being developed by a handful of companies operating with virtually no regulatory oversight or transparency, amplifying pre-existing societal inequalities while promising an unattainable utopia.

About the Author

Parmy Olson is a Bloomberg Opinion columnist who has covered the technology industry for over thirteen years. She previously worked as a reporter for the *Wall Street Journal* and *Forbes*. Olson is the author of *We Are Anonymous* and has received recognition for her business journalism. Her extensive background includes being the first to report on a secret effort at Google's top AI lab to break away from the company, a dramatic event that this book expands upon

Books Book Review Book Recommendations

Written by shebbar

82 followers · 28 following

Follow

No responses yet

Write a response

What are your thoughts?

More from shebbar

shebbar

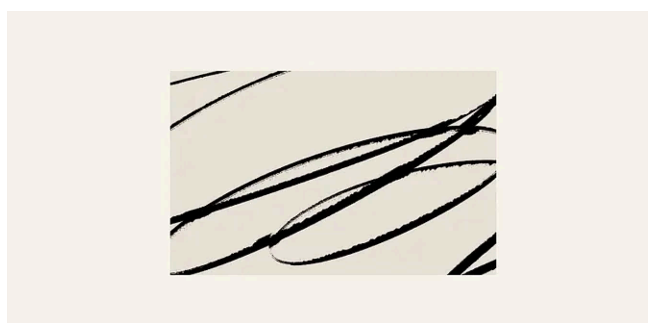

shebbar

AI Coding Agents & IDEs (36 Tools)

The Hidden Power of 'Glue Employees' and Quiet Leadership

In most organizations, success stories are often told through the lens of visible...

Oct 16, 2025 1

The future of coding is here, and it's powered by AI. From no-code platforms that turn...

Feb 25, 2025 3

shebbar

Federated Learning: The Future of Distributed Intelligence Through...

Executive Summary

Oct 20, 2025

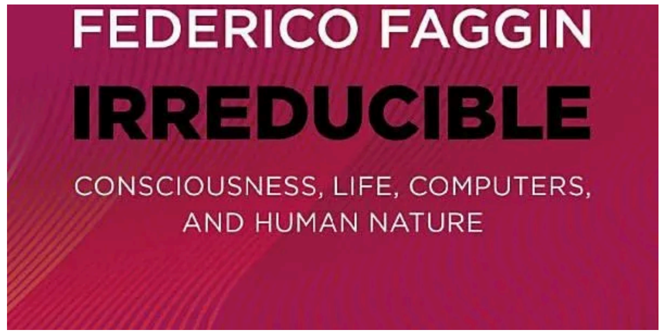

shebbar

A Journey Beyond the Material: Exploring Federico Faggin's...

Introduction:

Jul 15, 2024 7

See all from shebbar

Recommended from Medium

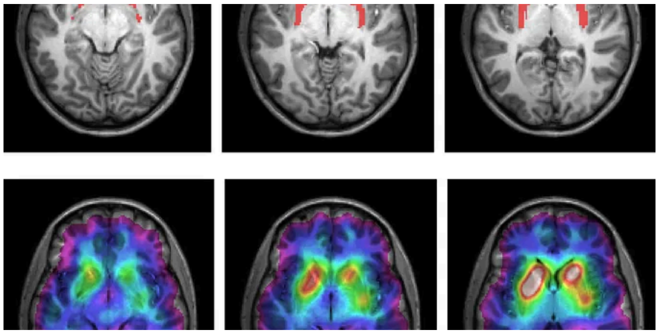

In Write A Catalyst by Dr. Patricia Schmidt

As a Neuroscientist, I Quit These 5 Morning Habits That Destroy You...

Most people do #1 within 10 minutes of waking (and it sabotages your entire day)

Jan 14 45K 917

In ILLUMINATION by Sufyan Maan, M.Eng

I Woke Up at 4:30 AM Every Day for 30 Days — Here Is What Nobody...

Here is what actually happened, from someone who did it & tracked everything.

4d ago 7K 310

In Code Like A Girl by Shruti Mangawa

Why Reading More Books Wasn't Making Me Smarter

I Realized I Had Been Reading Wrong My Entire Life

Dec 11, 2025 12K 502

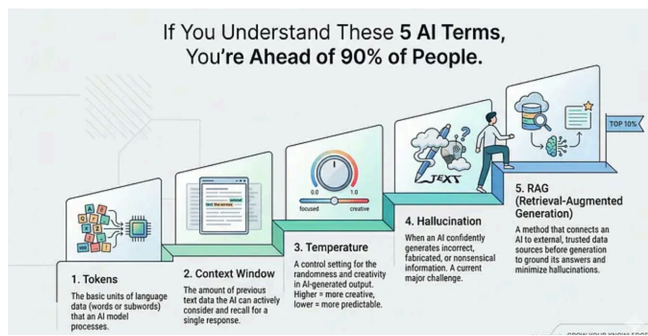

In Towards AI by Shreyas Naphad

If You Understand These 5 AI Terms, You're Ahead of 90% of...

Master the core ideas behind AI without getting lost

Mar 29 5.1K 111

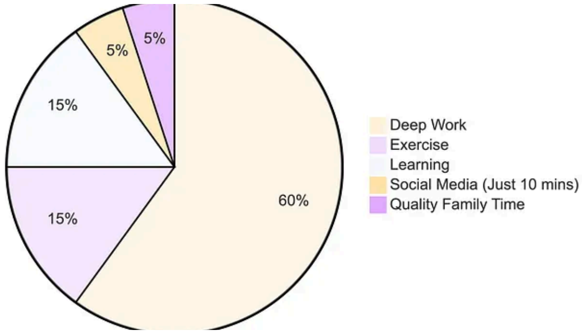

In Level Up Coding by Kusireddy

I Stopped Using ChatGPT for 30 Days. What Happened to My Brai...

91% of you will abandon 2026 resolutions by January 10th. Here's how to be in the 9% who...

Dec 28, 2025 12.6K 483

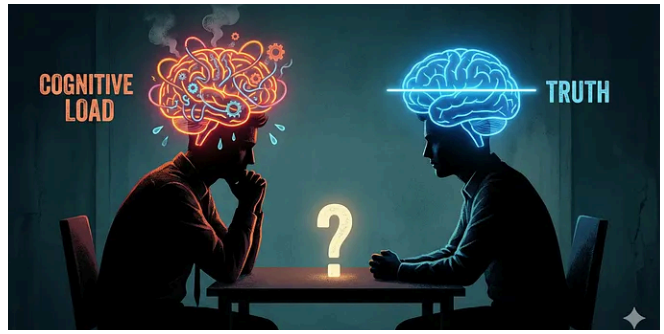

In E³—Entertain Enlighten Empo... by Joshua Mas...

How to Spot a Liar in 3 Questions (Without Them Knowing)

The secret to detecting deception isn't about body language—it's about "Cognitive Load."

Dec 11, 2025 23K 518

See more recommendations